



DELYAN BOYCHEV

@ delyan.boychev@insait.ai | [in d-boychev](#) | [github delyan-boychev](#) | [Scholar](#) | dboychev.com

Sofia, Bulgaria

ABOUT ME

Informatics undergraduate and INSAIT research scholar working on vision-language models, remote sensing, 3D scene understanding, and efficient transformer inference. IOAI lecturer; previously a Quant ML intern at Man Group and an NLP research intern at Graphwise. Research spans domain adaptation, synthetic-image detection, adversarial robustness, and interpretability.

EDUCATION

INSAIT Scholar (highly selective 4-year program)

Sofia, Bulgaria

Faculty of Mathematics and Informatics, Sofia University

Bachelor of Informatics

Sept. 2024 – Expected 2028

EXPERIENCE

NLP Research Intern | Graphwise

July 2026 – Sept. 2026

- Built NLP pipelines for table entity reconciliation, knowledge-base linking, and structured extraction from unstructured text.
- Optimized language-model pipelines for deployment, improving inference cost and latency.

IOAI Lecturer | Bulgarian IOAI Team

June 2025 – Present

- Lecturing and mentoring students preparing for the International Olympiad in Artificial Intelligence (IOAI).

Undergrad Research Scholar | INSAIT

Oct. 2024 – Present

- Researching vision-language models for remote sensing and 3D scene understanding under the advisory of Dr. Danda Paudel.

Quant ML Intern | Man Group

July 2025 – Sept. 2025

- Conducted Quant ML research, focusing on developing and evaluating machine-learning models for systematic trading and signal forecasting.

PUBLICATIONS

OSMDA: OpenStreetMap-based Domain Adaptation for Remote Sensing VLMs | *arXiv Preprint, 2026*

- Stefan Maria Ailuro, Mario Markov, Mohammad Mahdi, **Delyan Boychev**, Luc Van Gool, Danda Pani Paudel
- Self-contained domain adaptation using OpenStreetMap supervision without manual labels or a stronger external teacher; under review.
- **Paper:** arxiv.org/abs/2603.11804

ImagiNet: A Multi-Content Benchmark for Synthetic Image Detection | *AAAI 2025 Workshop Paper*

- **Delyan Boychev**, Radostin Cholakov — AAAI 2025 Workshop on Datasets and Evaluators of AI Safety
- A balanced 200K-image benchmark spanning four content types and multiple open and proprietary generators, with strong contrastive-learning baselines.
- **Paper:** arxiv.org/abs/2407.20020

Interpretable Computer Vision Models through Adversarial Training | *arXiv Preprint, 2023*

- Single-author empirical study of how adversarial training changes learned representations and model interpretability.
- **Paper:** arxiv.org/abs/2307.02500

PROJECTS

- DisentangledFlash** | *Triton Attention Kernel* Aug. 2026 – Present
- Fuses DeBERTa-v2/v3 attention into a streaming Triton kernel without materializing the full attention matrix.
 - On H200 benchmarks: up to $2.33\times$ faster and up to 91% lower peak memory than Hugging Face Transformers eager.
 - **Code:** github.com/delyan-boychev/disentangled-flash | **PyPI:** pypi.org/project/disentangled-flash
- LavBench** | *ML Competition Platform* June 2026 – Present
- A secure, bilingual platform for machine-learning competitions, built for Bulgarian AI Olympiad selection and national contests.
 - Runs notebook and Python submissions in hardened Docker sandboxes; provides live leaderboards, double-blind jury review, custom evaluators, and distributed workers.
 - **Code:** github.com/delyan-boychev/lavbench | **Docs:** lavbench.readthedocs.io
- Generative Translation** | *Coursework in Information Search and Retrieval: Application of Deep Learning* Feb. 2026
- Built a decoder-only PyTorch Transformer for Bulgarian-to-English translation with custom BPE, masked attention, KV caching, and beam search.
 - Achieved 42.57 BLEU-4 with 12.26 test perplexity using beam search.
 - **Code:** github.com/delyan-boychev/gen-trans
- Vision-Language Models for Satellite Imagery Understanding** | *Research at INSAIT* Dec. 2025 – Present
Advised by Dr. Danda Paudel
- Adapting vision-language models to satellite imagery using scalable supervision derived from OpenStreetMap.
- Spatial Modeling for Point-Cloud Understanding** | *Research at INSAIT* June 2025 – Present
Advised by Dr. Danda Paudel
- Aligning 3D point-cloud and Gaussian representations for structured spatial reasoning and scene understanding.
- C++ Coursework** | *DSA and OOP Course Projects* 2025
- Built a dependency-aware expression solver with polymorphic tokenization and two-stack evaluation, and a modular digital audio workstation with timeline editing, effects, and audio I/O.
 - **Solver:** [partial-solver](#) | **Audio Workstation:** [dig-audio-workstation](#)
- ImagiNet** | *Synthetic Image Dataset* 2024
- A 144 GB, 200K-image dataset for synthetic-image detection, balanced across photos, paintings, faces, and uncategorized imagery.
 - Supports both real-versus-synthetic classification and generator identification across open and proprietary image generators.
 - **Dataset:** huggingface.co/datasets/delyanboychev/imagenet
- Gradient Cache Contrastive Learning** | *PyTorch Package* 2023 – 2024
- A PyTorch implementation of gradient caching for scaling computer-vision contrastive learning beyond GPU memory limits.
 - Reproduces large-batch gradients from smaller chunks; supports SimCLR, SupCon, SelfCon, mixed precision, and distributed training.
 - **Code:** github.com/delyan-boychev/grad-cache-con-learning | **PyPI:** pypi.org/project/grad-cache-con-learning

HONORS

- Scholarship Recipient, EXPLORER Program, INSAIT (2024)
- European Space Agency Award from EUCYS'24, Katowice, Poland
- International Olympiad in Artificial Intelligence 2024, Team Bronze Medal
- Regeneron ISEF'23 Finalist, Dallas, Texas
- John Atanasoff Presidential Award in the category “Debut breakthrough in computer technology” (2023)

TECHNICAL SKILLS

- **Languages:** Python, C++, TypeScript
- **ML/Research:** PyTorch, Triton, vLLM, Hugging Face Transformers
- **Data Science:** NumPy, pandas, scikit-learn
- **Tools:** Linux, Docker, CI/CD, Git

LANGUAGES

- Bulgarian — Native
- English — IELTS Academic, 7.0/9.0